

# TBCX for Tcl 9.1: Design, Format, and Implementation Notes

TBCX Project

September 2025

## Abstract

This document describes the TBCX (Tcl ByteCode eXchange) file format and the accompanying `tbcx` extension for Tcl 9.1. TBCX allows precompiling Tcl scripts to a compact, portable, little-endian binary representation that can be loaded later without recompilation. We cover the public API, the exact on-disk format, loader behavior, and implementation details relevant to embedding and tooling.

## 1 Overview and Goals

TBCX goals:

- **Zero source recompilation at load time:** ship precompiled bodies.
- **Portability:** files are little-endian and validated at load.
- **Determinism:** dictionaries and jump tables are serialized in a stable order.
- **Safety:** only supported AuxData kinds are encoded; inputs are size-checked.

The package provides `tbcx::save`, `tbcx::savechan`, `tbcx::savefile`, `tbcx::loadchan`, `tbcx::loadfile`, and `tbcx::dumpfile`. The package version is 1.0.

## 2 Public API

### 2.1 Saving

`tbcx::save script outPath`

Compile *script* and write a TBCX file.

`tbcx::savechan script channelName`

Compile and write to an open channel; the command configures the channel for binary I/O.

`tbcx::savefile in.tcl outPath`

Read *in.tcl*, compile, and write a TBCX file.

### 2.2 Loading and Inspecting

`tbcx::loadfile in.tbcx`

Load and evaluate a TBCX file in the current namespace.

`tbcx::loadchan channelName`

Load and evaluate from a channel (configured for binary).

`tbcx::dumpfile filename`

Return a text dump (header, sections, disassembly) for inspection.

### 3 Stream Header

All integers are little-endian. The header layout is:

```
struct TbcxHeader {
    u32 magic;          // "TBCX" = 0x58434254
    u32 format;         // 9
    u32 tcl_version;    // maj<<24 | min<<16 | patch<<8 | type
    u64 codeLenTop;
    u32 numCmdsTop;     // always 0 in this version
    u32 numExceptTop;
    u32 numLitsTop;
    u32 numAuxTop;
    u32 numLocalsTop;
    u32 maxStackTop;
};
```

The loader validates *magic* and *format* before continuing.

### 4 Compiled Block Encoding

A compiled block serializes one `ByteCode` object in five parts:

1. **Code**: u32 byte length, then raw code bytes.
2. **Literals**: u32 count, then tagged literals (Section 5).
3. **AuxData**: u32 count, then tagged AuxData (Section 6).
4. **Exceptions**: u32 count, then range table (fields as u32).
5. **Epilogue**: u32 maxStack, u32 reserved (0), u32 numLocals.

For top-level compilation the default namespace is the current one at load time. For nested `lambdaExpr` and `bytecode` literals the stream precedes the block with a namespace FQN (LPString) and ensures that namespace exists before instantiation.

#### 4.1 Limits

Enforced maxima: code  $\leq$  1 GiB; literals  $\leq$  64 M; AuxData  $\leq$  64 M; exceptions  $\leq$  64 M; any LPString  $\leq$  16 MiB.

### 5 Literals

Each literal begins with a u32 tag:

- 0 **BIGNUM**: u8 sign (0=zero, 1=+, 2=-), u32 magLen, then magLen LE bytes.
- 1 **BOOLEAN**: u8 (0/1).
- 2 **BYTEARR**: u32 length + bytes.
- 3 **DICT**: u32 pair count, then key literal + value literal (keys sorted by UTF-8).
- 4 **DOUBLE**: 64-bit IEEE-754 as u64.
- 5 **LIST**: u32 element count, then that many nested literals.
- 6 **STRING**: LPString (u32 length + bytes).
- 7 **WIDEINT**: signed 64-bit two's complement (u64 payload).
- 8 **WIDEUINT**: unsigned 64-bit integer (promoted to bignum by loader if needed).
- 9 **LAMBDA\_BC**: ns FQN (LPString), u32 numArgs; for each arg: name (LPString), u8 hasDefault, op
- 10 **BYTECODE**: ns FQN (LPString) + compiled block.

**Lambda bodies.** For **LAMBDA\_BC** the saver marshals the public *args* list and compiles the body using proc semantics, creating a temporary **Proc** whose argument locals match the lambda signature.

## 6 AuxData

Supported AuxData tags:

### **JT\_STR (0)**

u32 count; for each entry: LPString key, u32 target offset.

### **JT\_NUM (1)**

u32 count; for each entry: u64 key, u32 target offset.

### **DICT\_UPDATE (2)**

u32 length; then that many local indices.

### **NEW\_FOREACH (3), FOREACH (4)**

u32 numLists, u32 loopCtTemp, u32 firstValueTemp, duplicate numLists; for each list:  
u32 numVars, then indices.

The loader maps either FOREACH tag to the closest available core AuxData type.

## 7 Procedures, Classes, and Methods

### 7.1 Static capture and sections

The saver parses the source text to capture static definitions: **proc name args body, namespace eval ns body**, and **oo::define** forms for methods/ctors/dtors, plus **oo::class create**. For each captured proc it stores three LPStrings (name as written, ns FQN, and the canonical args list) followed by the compiled body. Classes are written with their FQN and a list of superclasses (0 for now). Methods are written with class FQN, kind (inst/class/ctor/dtor), optional name, args list, and compiled body.

### 7.2 Load-time reconstruction

The loader installs temporary shims:

- **proc shim**: on **proc name args body**, if the fully-qualified name and *args* match an entry from the file, substitute the precompiled body; otherwise forward unchanged.
- **oo::define shim**: for **method**, **classmethod**, **constructor**, and **destructor**, substitute the precompiled body when the *args* list matches.

After reading all sections the loader evaluates the top-level block (with the current namespace as default), then restores the original commands and discards the shims.

### 7.3 Locals and variable frames

For each compiled body the saver records **numLocals**. The loader extends the **Proc**'s compiled-local list to at least that length to keep Tcl's invariants when freeing or redefining the proc later.

## 8 Disassembly Tool

The command `tbcx::dumpfile` produces a readable dump: header, compiled-block structure, AuxData, exceptions, and a disassembly. The decoder table is aligned to Tcl 9.1's `tclInstructionTable`. Opcodes marked “deprecated” correspond to legacy one-byte forms and may still be shown to aid debugging.

## 9 Portability and Limits

Streams are little-endian; both writer and reader swap on big-endian hosts. Size limits at save and load prevent pathological inputs (1 GiB code; pools up to 64 M elements; LPString up to 16 MiB).

## 10 Examples

### 10.1 Proc and class

```
set s {
  proc add {a b} { expr {$a + $b} }
  oo::class create C
  oo::define C method m {x} { expr {$x * 2} }
}
tbcx::save $s out.tbcx
tbcx::loadfile out.tbcx
puts [add 2 3] ;# => 5
puts [[C new] m 10] ;# => 20
```

### 10.2 Channels

```
set ch [open "bundle.tbcx" w]
fconfigure $ch -translation binary -eofchar {}
tbcx::savechan {proc p {} {return ok}} $ch
close $ch

set ch [open "bundle.tbcx" r]
fconfigure $ch -translation binary -eofchar {}
tbcx::loadchan $ch
close $ch
```

## 11 Build and Initialization Notes

The extension uses Tcl's stubs and TomMath stubs. During initialization the package resolves the necessary `Tcl_ObjTypes` (`bytecode`, `list`, `dict`, `int`, `double`, `boolean`, `bytearray`, `bignum`, and optionally `lambdaExpr`) and the AuxData type pointers (`JumptableInfo`, `JumptableNumInfo`, `DictUpdateInfo`, `ForeachInfo`, `NewForeachInfo`).

**Endianness.** Endianness is detected once (via `tcl_platform(byteOrder)` or a probe) and cached; streams are always written and read as little-endian.

## Acknowledgments

Thanks to the Tcl core for public/internal APIs enabling safe precompilation and to the reviewers who improved the serializer/loader design.

## 12 License

MIT License. Copyright © 2025 Miguel Bañón.